

Assignment 5: Ensemble Methods

Copyright and Fair Use

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

Automatic Testing Guidelines

Automatic unittesting requires you, as a student, to submit a notebook which contains strictly defined objects. Strictness of definition consists of unified shapes, dypes, variable names and more.

Within the notebook, we provide detailed instruction which you should follow in order to maximise your final grade.

Name your notebook properly, follow the pattern in template name:

Assignment_N_NameSurname_mnumber

1. N - number of assignment
2. NameSurname - your full name where every part of the name starts with a capital letter, no spaces
3. mnumber - your 8-digit student number on ID card (without k)

Example:

- Assignment_0_ReneDescartes_12345678
- Assignment_0_SemprahyKienegard_12345678
- Assignment0_Peter_Pan_12345678

Don't add any cells but use the ones provided by us. You may notice that most cells are tagged such that the unittest routine can recognise them.

We highly recommend you to develop your code within the provided cells. You can implement helper functions where needed unless you put them in the same cell they are actually called. Always make sure that implemented functions have the correct output and given variables contain the correct data type. Don't import any other packages than listed in the cell with the "imports" tag.

Note: Never use variables you defined in another cell in your functions directly, always pass them to the function as a parameter. In the unittest they won't be available either.

Good luck!

Task 1: AdaBoostM1 is an instance of forward stagewise modelling

In the lecture it was mentioned that one of the first boosting algorithms, i.e. AdaBoostM1, is equivalent to forward stagewise modelling using the exponential loss $L(y, g(\mathbf{x})) = \exp(-y g(\mathbf{x}))$ for a binary classification problem. In this task we need to provide proof of this fact. We will guide you through the most important steps and you will have to add some details.

For AdaBoostM1, the basis functions at timestep n are the individual classifiers $h_n(\mathbf{x}) \in \{-1, 1\}$. We assume that all of them are slightly better than random guessing. Note that we use h_n here for the resulting classifier at timestep n , which differs slightly from the notation in the slides, mainly to not confuse it with the corresponding approximation from forward stagewise modelling, which is also called g_n there.

Using the exponential loss in each timestep n we have to solve

$$(\beta_n, b_n) = \arg \min_{\beta, b} \sum_{i=1}^N \exp(-\beta (y_i - b (h_n(\mathbf{x}_i) + \beta(h(\mathbf{x}_i))),$$

for the classifier h_n and the coefficient β_n which are added at each step. This can be rewritten as

$$(\beta_n, b_n) = \arg \min_{\beta} \sum_{i=1}^N w_i^{(n)} \exp(-\beta (y_i - b(h_n(\mathbf{x}_i))), \quad (1)$$

with $w_i^{(n)} = \exp(-y_i g_{n-1}(\mathbf{x}_i))$. Since each $w_i^{(n)}$ depends neither on β nor b , it can be regarded as a weight that is applied to each observation. This weight depends on $g_{n-1}(\mathbf{x}_i)$, and so the individual weight values change with each iteration n . The solution of (1) can be found in two steps:

Calculation 1 (10 points):

In the first step fix $\beta \geq 0$ and show that in this case the solution to (1) is $b_n = \arg \min_{b \in \{-1, 1\}} \sum_{i=1}^N w_i^{(n)} I(h(\mathbf{x}_i) \neq y_i)$. Hints:

1. Try to write the expressions in (1) after argmin in the form $\sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} + \dots + \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}$.
- Find the right expressions for $I(\cdot)$, such that the product $b(h_n(\mathbf{x}_i))$ doesn't appear there anymore.

2. Now show that this can be written as

$$(\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)} \quad (2)$$

The equation $\sum_{i=1}^N w_i^{(n)} = \sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} + \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}$ might be helpful.

3. Argue why this already implies the claim.

Please provide reasoning and explanations in full sentences. Grading of the task will heavily depend on it.

Write your first step calculation here*

1. Hint 1 ...
2. Hint 2 ...
3. Hint 3 ...

$$1. (\exp(\beta), b_n) = \arg \min_{b \in \{-1, 1\}} \sum_{i=1}^N w_i^{(n)} \exp(-y_i b(h_n(\mathbf{x}_i)))$$

Splitting between datapoints correctly classified $y_i b(h_n(\mathbf{x}_i)) = 1$ and $y_i b(h_n(\mathbf{x}_i)) = -1$

$$(\beta_n, b_n) = \sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} \exp(-\beta) + \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} \exp(-\beta)$$

$$2. \text{Using } \sum_{i=1}^N w_i^{(n)} = \sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} + \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}$$

$$\Rightarrow \sum_{h(\mathbf{x}_i)=y_i} w_i^{(n)} = \sum_{i=1}^N w_i^{(n)} - \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}$$

$$(\beta_n, b_n) = \sum_{i=1}^N w_i^{(n)} \exp(-\beta) - \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} \exp(-\beta) + \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} \exp(-\beta)$$

$$= \sum_{i=1}^N w_i^{(n)} \exp(-\beta) + (\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}$$

$$= (\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

$$3. \text{The expression from (1) can be rewritten to } b_n = \arg \min_{b \in \{-1, 1\}} \sum_{i=1}^N w_i^{(n)} I(h_n(\mathbf{x}_i) \neq y_i)$$

because it can be minimized only if the sum of the weights is minimized, hence it makes sense to rewrite it with the identity matrix.

Calculation 2 (20 points):

In the second step you need to optimize the following expression with respect to β :

$$(\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

Hint: do it in the usual way (i.e. by differentiating the expression and setting it to 0).

Using the abbreviation $\text{err}_n = \frac{\sum_{i=1}^N w_i^{(n)} I(y_i \neq h_n(\mathbf{x}_i))}{\sum_{i=1}^N w_i^{(n)}}$ show that the obtained expression is $\beta_n = \frac{1}{2} \ln \frac{1 + \text{err}_n}{1 - \text{err}_n}$.

Note that $\beta_n \geq 0$, by our assumption that all classifiers are better than random guessing, i.e. $\text{err}_n \leq \frac{1}{2}$, so the result is in accordance with the previous subtask.

Please provide reasoning and explanations in full sentences. Grading of the task will heavily depend on it.

Write your second step calculation here*

1. Derivate ...
2. Using given abbreviation show ...

$$1. (\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

$$= \exp(\beta) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} - \exp(-\beta) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

$$\Rightarrow \frac{\partial}{\partial \beta} \left((\exp(\beta) - \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{i=1}^N w_i^{(n)} \right)$$

$$= \exp(\beta) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} + \exp(-\beta) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} - \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

$$= (\exp(\beta) + \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} - \exp(-\beta) \sum_{i=1}^N w_i^{(n)}$$

$$2. (\exp(\beta) + \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} - \exp(-\beta) \sum_{i=1}^N w_i^{(n)} = 0$$

$$\Leftrightarrow (\exp(\beta) + \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} = \exp(-\beta) \sum_{i=1}^N w_i^{(n)} \mid \cdot \ln$$

$$\Leftrightarrow \ln((\exp(\beta) + \exp(-\beta)) \sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}) = \ln(\exp(-\beta) \sum_{i=1}^N w_i^{(n)})$$

$$\Leftrightarrow \ln((\exp(\beta) + \exp(-\beta))) + \ln\left(\sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)}\right) = \ln(\exp(-\beta)) + \ln\left(\sum_{i=1}^N w_i^{(n)}\right)$$

$$\Leftrightarrow \ln((\exp(\beta) + \exp(-\beta))) + \ln\left(\sum_{h(\mathbf{x}_i) \neq y_i} w_i^{(n)} I(h_n(\mathbf{x}_i) \neq y_i)\right) = -\beta + \ln\left(\sum_{i=1}^N w_i^{(n)}\right)$$

$$\Leftrightarrow \ln((\exp(\beta) + \exp(-\beta))) + \ln(\exp(\beta)) = \ln\left(\sum_{i=1}^N w_i^{(n)}\right) - \ln\left(\sum_{i=1}^N w_i^{(n)} I(h_n(\mathbf{x}_i) \neq y_i)\right)$$

$$\text{Since } \text{err}_n = \frac{\sum_{i=1}^N w_i^{(n)} I(y_i \neq h_n(\mathbf{x}_i))}{\sum_{i=1}^N w_i^{(n)}}$$

$$\Rightarrow \frac{1}{\text{err}_n} = \frac{\sum_{i=1}^N w_i^{(n)}}{\sum_{i=1}^N w_i^{(n)} I(y_i \neq h_n(\mathbf{x}_i))}$$

$$\Rightarrow \ln((\exp(\beta)^2 + \exp(-\beta) \exp(\beta))) = \ln\left(\frac{1}{\text{err}_n}\right)$$

$$\Rightarrow \ln(\exp(\beta)^2 + 1) = \ln\left(\frac{1}{\text{err}_n}\right)$$

$$\Rightarrow \exp(\beta)^2 + 1 = \frac{1}{\text{err}_n}$$

$$\Rightarrow \exp(\beta)^2 = \frac{1 - \text{err}_n}{\text{err}_n} \mid \cdot \ln$$

$$\Rightarrow 2\beta = \ln\left(\frac{1 - \text{err}_n}{\text{err}_n}\right)$$

$$\Rightarrow \beta_n = \frac{1}{2} \ln\left(\frac{1 - \text{err}_n}{\text{err}_n}\right)$$

Calculation 3 (10 points):

In the final step we can update the approximation as follows: $g_n(\mathbf{x}) = g_{n-1}(\mathbf{x}) + \beta_n h_n(\mathbf{x})$. To finish the proof proceed by deriving the following relations:

1. The weights for the next generation can be computed as follows: $w_i^{(n+1)} = w_i^{(n)} \exp(-y_i \beta_n h_n(\mathbf{x}_i))$.
2. $-y_i h_n(\mathbf{x}_i) - 2I(y_i \neq h_n(\mathbf{x}_i)) - 1$
3. Use these two relations to show that:

$$w_i^{(n+1)} = w_i^{(n)} \exp(-\beta_n) \exp(\alpha_n I(y_i \neq h_n(\mathbf{x}_i))) \quad (2)$$

where $\alpha_n = -2\beta_n$ is the α_n from the AdaBoostM1 algorithm from the lecture.

Please provide reasoning and explanations in full sentences. Grading of the task will heavily depend on it.

Write your final step calculation here*

1. Calculate weights ...
2. ...
3. ...

$$1. \text{As stated in the task description: } w_i^{(n+1)} = \exp(-y_i g_n(\mathbf{x}_i))$$

$$\Rightarrow w_i^{(n+1)} = \exp(-y_i g_n(\mathbf{x}_i))$$

$$\text{Using } g_n(\mathbf{x}) = g_{n-1}(\mathbf{x}) + \beta_n h_n(\mathbf{x}) \Rightarrow$$

$$= \exp(-y_i g_{n-1}(\mathbf{x}_i) - y_i \beta_n h_n(\mathbf{x}_i))$$

$$= w_i^{(n)} \exp(-y_i \beta_n h_n(\mathbf{x}_i))$$

2. There are two possibilities: either the datapoint is correctly classified, either misclassified

$$\text{I. If datapoint is correctly classified} \Rightarrow y_i h_n(\mathbf{x}_i) = 1 \mid \cdot (-1)$$

$$\Leftrightarrow -y_i h_n(\mathbf{x}_i) = -1$$

$$\text{If } y_i h_n(\mathbf{x}_i) = 1 \Rightarrow I(y_i \neq h_n(\mathbf{x}_i)) = 0 = 2I(y_i \neq h_n(\mathbf{x}_i))$$

$$\Rightarrow -y_i h_n(\mathbf{x}_i) = 2I(y_i \neq h_n(\mathbf{x}_i)) - 1 \quad (q.e.d.)$$

$$\text{II. If datapoint is misclassified} \Rightarrow y_i h_n(\mathbf{x}_i) = -1 \mid \cdot (-1)$$

$$\Leftrightarrow -y_i h_n(\mathbf{x}_i) = 1$$

$$\text{If } y_i h_n(\mathbf{x}_i) = -1 \Rightarrow I(y_i \neq h_n(\mathbf{x}_i)) = 1 \Rightarrow 2I(y_i \neq h_n(\mathbf{x}_i)) = 2$$

$$\Rightarrow -y_i h_n(\mathbf{x}_i) = 2I(y_i \neq h_n(\mathbf{x}_i)) - 1 \quad (q.e.d.)$$

Thus $-y_i h_n(\mathbf{x}_i) = 2I(y_i \neq h_n(\mathbf{x}_i)) - 1$ holds even though the datapoints are classified correctly or misclassified

$$3. \text{Using the formula from subtask 1: } w_i^{(n+1)} = w_i^{(n)} \exp(-y_i \beta_n h_n(\mathbf{x}_i))$$

$$\text{Using the formula from subtask 2: } -y_i h_n(\mathbf{x}_i) = 2I(y_i \neq h_n(\mathbf{x}_i)) - 1 \Rightarrow$$

$$w_i^{(n+1)} = w_i^{(n)} \exp(\beta_n (2I(y_i \neq h_n(\mathbf{x}_i)) - 1))$$

$$= w_i^{(n)} \exp(2\beta_n I(y_i \neq h_n(\mathbf{x}_i)) - \beta_n)$$

$$= w_i^{(n)} \exp(-\beta_n) \exp(2\beta_n I(y_i \neq h_n(\mathbf{x}_i)))$$

$$\text{with } 2\beta_n = \alpha_n \Rightarrow w_i^{(n+1)} = w_i^{(n)} \exp(-\beta_n) \exp(\alpha_n I(y_i \neq h_n(\mathbf{x}_i)))$$

Task 2: Random Forests and Feature importance

In this task you will train a Random Forest (RF) Classifier on a subset of fashionMNIST. You should observe how these models can immediately give you useful information about feature importance, which is a very convenient property of RFs.

First you should re-use the given code from the previous assignment to load the whole data set (the procedure is completely analogous).

- Next implement code that filters the data set for the classes with labels 1 (trousers) and 9 (ankle boots). You should create the filtered train data set from `x_train` and `y_train` and the test data set from `x_test` and `y_test`.

- Hint: Masks provide a convenient solution to this task.
- After the filtering procedure the data samples corresponding to trousers should be labelled as 1 and the ankle boots as 0. Perform this step on the test and train data set.
- To accomplish this task, implement a function `_filter_0`.

Code 1 (15 points):

```
In [1]: """NOTE: Please add all your imports in this cell only"""
#Please add all your imports in this cell only
#Imports
#nothing to do here
import numpy as np
import pandas as pd
import sys
import time
import sklearn.ensemble import RandomForestClassifier
from miscl_loader import MISCL
import matplotlib.pyplot as plt
from matplotlib import style
import matplotlib as mpl
import seaborn as sns
from sklearn.metrics import accuracy_score, confusion_matrix
# set random seed to ensure reproducible runs
RSEED = 10
```

```
In [2]: #Load training and test data (routine from last week)
data = MISCL('./dataset/')
img_train, labels_train = data.load_training()
x_train = np.array(img_train)
y_train = np.array(labels_train)
x_test, y_test = data.load_testing()
x_test = np.array(x_test)
y_test = np.array(y_test)
print(y_train)
print(y_test)
```

```
In [3]: """Function _filter_() is created to filter datasets w.r.t to given labels
@Param x_train Training Feature matrix
@Param y_train Training Labels vector
@Param x_test Test Feature matrix
@Param y_test Test Labels vector
@Param labels_list list of length 2 which consists of integer labels.
@Returns a tuple (x_train_filtered, y_train_filtered, x_test_filtered, y_test_filtered).
```

```
def _filter_(x_train:np.array,
            y_train:np.array,
            x_test:np.array,
            y_test:np.array,
            labels_list):
    #your code goes here ...
    labels1=labels_list[0]
    labels2=labels_list[1]
```

```
    x_train_filtered = x_train[np.where(np.logical_or(y_train==labels1,y_train==labels2))]
    y_train_filtered = y_train[np.where(np.logical_or(y_train==labels1,y_train==labels2))]
    x_test_filtered = x_test[np.where(np.logical_or(y_test==labels1,y_test==labels2))]
    y_test_filtered = y_test[np.where(np.logical_or(y_test==labels1,y_test==labels2))]
```

```
    x_train_filtered,x_train_filtered=0
    y_train_filtered,y_train_filtered=0
    x_test_filtered,y_test_filtered=0
    y_test_filtered,y_test_filtered=0
```


```
    #Your code ends here
    return (x_train_filtered, y_train_filtered, x_test_filtered, y_test_filtered)
```

```
In [4]: #NOTHING TO CHANGE HERE
x_train,y_train,x_test,y_test = _filter_(x_train,y_train,x_test,y_test,[1,9])
print(y_train)
print(y_test)
```

```
[0 0 ... 0 0 7]
[0 1 ... 0 0 1]
```

In the following we visualize a few randomly selected samples from our training data:

```
In [5]: #A routine that you can use for plotting some of the data.
def plot_images(images, titles, n_rows=10, n_cols=10, figsize=(10,20)):
    a = np.random.randint(1,40,20)
    plt.figure(figsize=(20,20))
    for n,i in enumerate(a):
        plt.subplot(2, n_cols, i)
        two_d = np.reshape(img_train[i], (28, 28)) * 255.0, dtype=np.uint8)
        plt.title('Label: %0' % format(arr[y_train[i]]))
        plt.imshow(two_d, interpolation='nearest', cmap='gray')
    plt.subplots_adjust(hspace = 0.3)
```



Code (5 points):

Your task now is to train a `sklearn.RandomForestClassifier` with the default parameters on the training data set.

Then get the model's predictions for the test data set. Use `RSEED` as `random_seed` for the `RandomForestClassifier`.

For this, we ask you to implement a function `fit_predict`.

```
In [6]: """Function fit_predict() is created to fit RF on training data and return predictions as well as model
@Param x_train Training Feature matrix
@Param y_train Training Labels vector
@Param x_test Test Feature matrix
@Param y_test Test Labels vector
@Param random_seed, integer
@Returns a tuple (model, prediction), where model is an trained RF classifier and prediction is a np array
def fit_predict(x_train,y_train,x_test,y_test, rseed):
    #your code goes here ...
    model = RandomForestClassifier(random_state=rseed)
    model.fit(x_train,y_train)
    prediction = model.predict(x_test)
    #Your code ends here
    return model, prediction
```

```
In [7]: model, prediction = fit_predict(x_train,y_train,x_test,y_test, RSEED)
```

Code (15 points):

Now, within variables `size_test`, `num_wrong` save the size of the test set and number of misclassified test samples. Both variables should be integers. Retrieve misclassified samples from test set and save them as a list `items_wrong`.

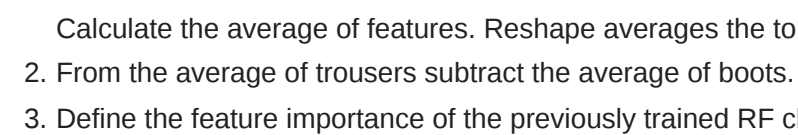
Use your predictions to plot up to 20 test data sample(s) that were misclassified:

```
In [8]: #NUMBER OF WRONG PREDICTIONS
#your code goes here ...
size_test = len(x_test)
num_wrong = 0
items_wrong = []
for value in range(len(prediction)):
    if prediction[value] != y_test[value]:
        num_wrong+=1
        items_wrong.append(x_test[value])
#Your code ends here
```

```
#Following print statement might be evaluated
print('Number of test samples: %d'%num_wrong)
print('Number of misclassified samples: %d'%size_test*num_wrong)
```

```
Number of test samples: 2000
Number of misclassified samples: 1
```

```
In [9]: #PLOTING WRONG PREDICTIONS
#your code goes here ...
#use your favorite plotting, or routine we have shown above
up to 20 samples from (20)
for i in up to 20 samples:
    break
fig=plt.figure()
two_d = np.reshape(items_wrong[i], (28, 28)) * 255.0, dtype=np.uint8)
plt.imshow(two_d, interpolation='nearest', cmap='gray')
```



Code (20 points):

Set within this part we will try to see the decision-making incentives of Random Forest.

To do this we ask you to implement the following tasks:

- Take your training dataset and split it into 2: trousers and boots.
- Calculate the average of features. Reshape averages to the 2D arrays of shape 28x28 and plot them as heatmaps. Save results under variables `tr_av` and `bo_av`.
- From the average of trousers subtract the average of boots. Save result as a variable `diff`. Plot it as a heatmap.
- Define the feature importance of the previously trained RF classifier with variable `importances`. Visualize it as a heatmap.

Hint: Check scikit-learn documentation to access feature importance.

The evaluation of the following code will be done by viewing your plots.

For those who are curious: run RF under different seeds, and look how plots are changing.

BEFORE SUBMISSION RETURN TO ORIGINAL SEED = 10


```
In [10]: #PLOTING HEATMAPS
#your code goes here ...
#use your favorite plotting
#no plots/variables are needed
#Step 1 Split and calculate averages
x_train_tr,x_train_bo=[]
for i in range(len(x_train)):
    if x_train[i]==1:
        x_train_tr.append(x_train[i])
    else:
        x_train_bo.append(x_train[i])
x_train_tr=np.array(x_train_tr)
x_train_bo=np.array(x_train_bo)
tr_av = np.average(x_train_tr,axis=0)
bo_av = np.average(x_train_bo,axis=0)
tr_av=np.reshape(tr_av,(28,28))
bo_av=np.reshape(bo_av,(28,28))
#Step 2 Subtract
temping_array=(tr_av,bo_av)
diff=tr_av-bo_av
```

```
#Step 3 Extract Feature Importance
importances = model.feature_importances_
importances=np.reshape(importances,(28,28))
```


```
#Step 4 plot everything together
fig=plt.figure(figsize=(12,2),figsize=(8,8))
sns.heatmap(data=tr_av,ax=axis(0,0))
axis(0,0).set_title('Average trousers')
sns.heatmap(data=bo_av,ax=axis(0,1))
axis(0,1).set_title('Average boots')
sns.heatmap(data=diff,ax=axis(1,0))
axis(1,0).set_title('Overlaps')
sns.heatmap(data=importances,ax=axis(1,1))
axis(1,1).set_title('Feature importances')
```

```
plt.show()
#Your code ends here
```

Main task



Main Task



If you have solved the previous task correctly, the resulting plot should look close to this:

Questions (5 points):

What observations can you make?

(Multiple answers might be correct)

To answer the question assign to variables in the next cell `True` or `False` boolean values. To earn points assign values to all variables.

Note: Do not reuse these variable names. They are used for testing.

a.) RF achieves accuracy lower than 85%